

POSIX

waitpid - 3 params

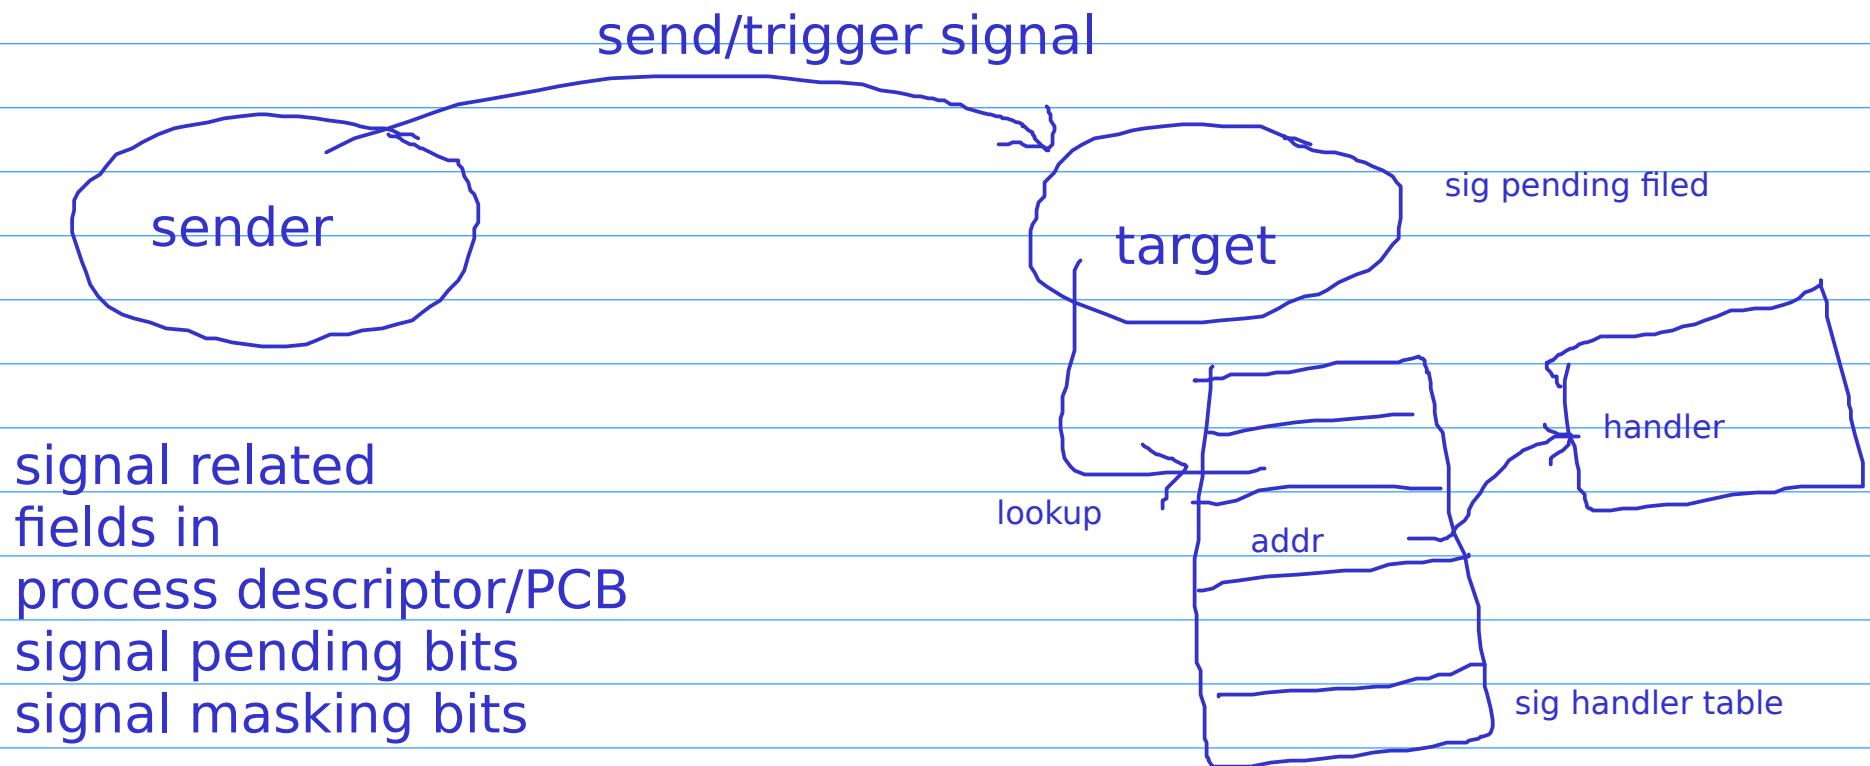
1st param : pid of child

2nd param : exit status (by address, fetch param)

3rd param : flags

execl, execlp, execv

Signals:-



when signal reaches target process, corresponding pending bit will be set

Most of default signal handlers will cause abnormal termination (typically)

ctrl + C	==>	SIGINT
ctrl + \	==>	SIGQUIT
ctrl + Z	==>	SIGTSTP
kill <pid>	==>	SIGTERM
child exit	==>	SIGCHLD
fg/bg	==>	SIGCONT
floating point	==>	SIGFPE
seg fault	==>	SIGSEGV

other:- SIGKILL, SIGSTOP, SIGPIPE, SIGALRM

kill -l

kill -SIGxxx <pid>

kill -<signo> <pid>

signal handling -- custom handler

Simple APIs:- signal, raise, pause, kill, alarm (Some are deprecated)

Modern:- sigaction, sigprocmask, sigsuspend

Non Maskable Signals:- SIGKILL, SIGSTOP (No custom handlers)

kill <pid>

kill -9 <pid>

kill(pid, signo);

kill(0, signo); //signal to self, like "raise"

Further explore:-

sigprocmask

sigalrm

sigaction

sigsuspend

killall, pkill, pgrep

fg, bg, jobs

command &

return 0 ==> exit(0) ==> _exit(0)
lib call wrapper for exit system call

What's diff b'n exit and _exit

```
for(i=1;i<=10;i++)  
    printf("hello:%d\t",i); // no \n  
exit / _exit (or) any exception before exit
```

Threads:-

- * part of application, corresponds sub activity
- * light weight process (part of process??)
- * flow of control

Significant:-

- * Concurrent execution (multiplexing of CPU)
- * Resource sharing
- * every process also runs as single thread initially

resource usage point of view ==> process
execution point of view ==> threads

multithreaded application

thread creation is lighter/faster than fork (LWP)

Need for threading/examples:-

- * Office Suite
- * Media Player
- * Browser
- * Concurrent server -- each thread for every client

task driven parallelism

data driven parallelism, e.g. sum of large array

some threads are CPU centric (no blocking),
some are I/O centric (freq block)

any access other CPU & memory ==> blocking call

threads will share all resources, except stack
every thread will maintain private stack

stack & context (PC) will vary from thread to thread

Self Study:-

Thread models

user level threads (many to one mapping)

kernel supported threads (one to one mapping)

Thread pools

Thread cancellation policies

Thread contention scope

CPU Affinity

`ps -e -L -o pid,ppid,lwp,nlwp,cmd`

POSIX Thread Library - pthread APIs

pthread_create (based on clone system call)

pthread_join

pthread_self

pthread_equal

pthread_yield

pthread_cancel

```
void* task_body(void* pv) {  
    printf("A--welcome\n");  
    //for loop  
}
```

psample.c

```
pthread_t pt1;  
pthread_create(&pt1, NULL, task_body, NULL);  
//....  
pthread_join(pt1, NULL);
```

pthread_create params:

- 1st : addr of pthread_t variable
- 2nd : attributes, NULL means default
- 3rd : entry point for execution
- 4th : arguments to entry function

```
gcc psample.c -lpthread  
-lpthread ==> libpthread.a libpthread.so
```

pthread_exit ==> current thread only, resources are not freed up
exit ==> process exit, resources are released
==> all threads will be terminated

Scheduling:-

Self explore:-

- * Scheduling params - CPU utilization, throughput
Turnaround time
scheduling latency/delay/jitter
response time
 - * Fairness among processes/threads
 - * No Starvation
 - * CPU Bound vs I/O Bound processes
 - * Preemptive vs Non preemptive scheduling
 - * Classical Algorithms
 - * First Come First Serves (FCFS/FIFO)
 - * Priority based scheduling
 - * Round Robin
 - * Shortest Job First (** Theoretical, not practical)
 - * Combination of algorithms, e.g. FIFO + RR, FIFO + RR
 - * Multi level scheduling, Multi level feedback queue scheduling
 - * Scheduling in SMP
-
- * Practical scheduling in Linux (very imp)

Who will be in ready state (ready queue)?


Real time -- deterministic
hard realtime vs soft realtime

Scheduling in Linux:-

- * Realtime policies
 - * SCHED_FIFO (PRIO + FIFO), 100 levels of priority
 - * SCHED_RR (PRIO + RR), 100 levels of priority
- * Nonrealtime policies
 - * SCHED_OTHER : Timesharing algorithm

Timesharing policy:

- * Modified round robin algos
- * 40 priority levels, controlled by nice vals
- * Time quantum is proportional to priority
- * Process may move among priority levels
 - * by explicit command / system calls
 - * based on CPU usage or waiting time [e.g. aging technique]
- * There are two queues in each priority levels
 - * Active & Expired
 - * When active queue is empty, swap the queues



```
ps -e -o pid,ppid,stat,policy,pri,ni,cmd
```

```
policy - TS, priority - 19, nice - 0
```

```
nice -n 5 ./a.out          # bc or ping
nice -n -6 ./a.out         # sudo
```

nice	prio	nice value range: -20 to +19
19	0	pri scale : 0 to 39
5	14	
0	19	
-6	25	
-20	39	

Activity:-

- * Assignment
- * IPC Pre-reading (Semaphores, Mutex, Message Queues, Shared Mem, Pipe/Fifos)
- * Virtual Memory (videos)