

Summary:-

In x86

Paging

32 bit addr field

Page Tables, Page Table Entries

TLB Cache

20 bit page no. +

Swap Area (Backing store for pages)

12 bit offset

Virtual Memory

10 bit primary +

Page Fault, Page Fault Exception Handler

Lazy Allocation/Demand Paging

10 bit secondary +

Thrashing

12 bit offset

Hierarchical Page Tables

Locality Principle

Page out

Major, Minor Faults

Page In

Copy on Write

Dirty Bit

Swap out

Page Locking (Pre-Fetching)

Swap In

Page Mapping -- Anonymous vs File

-- Private vs Shared Mapping

Linux Memory Zones:-

ZONE_DMA		Must be contiguous
ZONE_NORMAL	(LOW Memory)	May not be contiguous But can't be in swap area
ZONE_HIGH	(HIGHMEM)	Not contiguous, can be paged out (page faults)

e.g. 4 GB Memory
Around 896MB ==> LOWMEM (Around 25%)
Rest of mem ==> HIGHMEM (Around 75%)

Userspace memory allocation:- mmap system call:-

- Anonymous + Private
- Anonymous + Shared
- File Mapping + Private
- File Mapping + Shared

File Mapping ==> File contents are mapped to virtual pages
==> Any read/write to virtual pages will reflect
to actual file on disk (post sync up for write)

Anonymous Mapping ==> Blank Virtual Pages

MAP_ANONYMOUS
MAP_PRIVATE
MAP_SHARED

Hands-on:-

example3a.c	==>	file + private
example3b.c	==>	anonymous + private
example2.c	==>	shared + anonymous, between parent & child
example1a.c	==>	anonymous + private, observe vsz, min faults
example1b.c	==>	file + private (Let's skip for now)

example3a - fix:-

PROT_WRITE|PROT_READ

file should have more than 26 chars

example3b - fix

PROT_WRITE|PROT_READ

In case of shared memory, during fork, PTEs will be duplicated , but both PTEs will point to same frames (actual sharing)

Shared Memory b'n unrelated process (no parent-child)

shm_open

shm_close

shm_unlink

Analysis:-

shared memory

- * userspace buffer

- + no system calls required, simple memory operations (faster)

- no synchronization support (race conds may occur, sequencing may not be enforced)

message queues & pipes

- + kernel space buffer

- # send/write, recv/read are system calls

- + sync support is there , e.g. recv can block if Q is empty
send can append

message queues vs pipes

MQ - discrete messages (no merging/splitting)

Pipes - continous stream of data (data merges onw write, split on read)