

[NAME](#) | [SYNOPSIS](#) | [DESCRIPTION](#) | [RETURN VALUE](#) |  
[NOTES](#) | [EXAMPLE](#) | [SEE ALSO](#) | [COLOPHON](#)

  
Search online pages

## SYSCALL(2)

## Linux Programmer's Manual

## SYSC

### NAME [top](#)

`syscall` - indirect system call

### SYNOPSIS [top](#)

```
#define _GNU_SOURCE          /* See feature_test_macros(7) */
#include <unistd.h>
#include <sys/syscall.h>    /* For SYS_xxx definitions */

long syscall(long number, ...);
```

### DESCRIPTION [top](#)

`syscall()` is a small library function that invokes the system call whose assembly language interface has the specified *number* with specified arguments. Employing `syscall()` is useful, for example, when invoking a system call that has no wrapper function in the library.

`syscall()` saves CPU registers before making the system call, restores the registers upon return from the system call, and stores any error code returned by the system call in `errno(3)` if an error occurs.

Symbolic constants for system call numbers can be found in the header file `<sys/syscall.h>`.

### RETURN VALUE [top](#)

The return value is defined by the system call being invoked. In general, a 0 return value indicates success. A -1 return value indicates an error, and an error code is stored in `errno`.

### NOTES [top](#)

`syscall()` first appeared in 4BSD.

#### Architecture-specific requirements

Each architecture ABI has its own requirements on how system calls

arguments are passed to the kernel. For system calls that have glibc wrapper (e.g., most system calls), glibc handles the details of copying arguments to the right registers in a manner suitable for the architecture. However, when using **syscall()** to make a system call, the caller might need to handle architecture-dependent details; this requirement is most commonly encountered on certain 32-bit architectures.

For example, on the ARM architecture Embedded ABI (EABI), a 64-bit value (e.g., *long long*) must be aligned to an even register pair. Thus, using **syscall()** instead of the wrapper provided by glibc, **readahead()** system call would be invoked as follows on the ARM architecture with the EABI:

```
syscall(SYS_readahead, fd, 0,
        (unsigned int) (offset >> 32),
        (unsigned int) (offset & 0xFFFFFFFF),
        count);
```

Since the offset argument is 64 bits, and the first argument (*fd*) is passed in *r0*, the caller must manually split and align the 64-bit value so that it is passed in the *r2/r3* register pair. That means inserting a dummy value into *r1* (the second argument of 0).

Similar issues can occur on MIPS with the O32 ABI, on PowerPC with the 32-bit ABI, and on Xtensa.

The affected system calls are **fcntl64(2)**, **ftruncate64(2)**, **posix\_fadvise(2)**, **pread64(2)**, **pwrite64(2)**, **readahead(2)**, **sync\_file\_range(2)**, and **truncate64(2)**.

## Architecture calling conventions

Every architecture has its own way of invoking and passing arguments to the kernel. The details for various architectures are listed in the two tables below.

The first table lists the instruction used to transition to kernel mode, (which might not be the fastest or best way to transition to the kernel, so you might have to refer to **vdso(7)**), the register used to indicate the system call number, and the register used to return the system call result.

arch/ABI	instruction	syscall #	retval	Notes
arm/OABI	swi NR	-	a1	NR is syscall number
arm/EABI	swi 0x0	r7	r0	
blackfin	excpt 0x0	P0	R0	
i386	int \$0x80	eax	eax	
ia64	break 0x100000	r15	r10/r8	bool error/ errno value
parisc	ble 0x100(%sr2, %r0)	r20	r28	
s390	svc 0	r1	r2	See below
s390x	svc 0	r1	r2	See below
sparc/32	t 0x10	g1	o0	

sparc/64	t 0x6d	g1	o0
x86_64	syscall	rax	rax

For s390 and s390x, NR (the system call number) may be passed directly with "svc NR" if it is less than 256.

The second table shows the registers used to pass the system call arguments.

arch/ABI	arg1	arg2	arg3	arg4	arg5	arg6	arg7
arm/OABI	a1	a2	a3	a4	v1	v2	v3
arm/EABI	r0	r1	r2	r3	r4	r5	r6
blackfin	R0	R1	R2	R3	R4	R5	-
i386	ebx	ecx	edx	esi	edi	ebp	-
ia64	out0	out1	out2	out3	out4	out5	-
parisc	r26	r25	r24	r23	r22	r21	-
s390	r2	r3	r4	r5	r6	r7	-
s390x	r2	r3	r4	r5	r6	r7	-
sparc/32	o0	o1	o2	o3	o4	o5	-
sparc/64	o0	o1	o2	o3	o4	o5	-
x86_64	rdi	rsi	rdx	r10	r8	r9	-

Note that these tables don't cover the entire calling convention; architectures may indiscriminately clobber other registers not here.

## EXAMPLE [top](#)

```
#define _GNU_SOURCE
#include <unistd.h>
#include <sys/syscall.h>
#include <sys/types.h>
#include <signal.h>

int
main(int argc, char *argv[])
{
    pid_t tid;

    tid = syscall(SYS_gettid);
    tid = syscall(SYS_tgkill, getpid(), tid, SIGHUP);
}
```

## SEE ALSO [top](#)

[\\_syscall\(2\)](#), [intro\(2\)](#), [syscalls\(2\)](#), [vdso\(7\)](#)

## COLOPHON [top](#)

This page is part of release 3.70 of the Linux *man-pages* project description of the project, information about reporting bugs, and latest version of this page, can be found at <http://www.kernel.org/doc/man-pages/>.

**Linux****2014-05-10****SYSCA**

---

### Copyright and license for this manual page

---

HTML rendering created 2014-07-09 by [Michael Kerrisk](#), author of *The Linux Programming Interface*, maintainer of the [Linux man-pages](#) project.

For details of in-depth **Linux/UNIX system programming training courses** that I teach, look [here](#).

Hosting by [jambit GmbH](#).

