

Reparenting -- why to avoid (waitpid usage)

execl

./a.out 10 20 abc 2.3 xyz

```
int s1;
```

```
waitpid(pid, &s1, 0);
```

```
WIFEXITED(s1), WEXITSTATUS(s1)
```

```
execl("./a.out", "a.out", "10",  
"20", "abc", "2.3", "xyz", NULL);
```

```
ret=fork();
```

```
if(ret==0)
```

```
    //child code
```

```
else
```

```
    //parent code
```

```
char* argv[5] = {"a.out", "10",  
"20", "abc", "2.3", "xyz", NULL}
```

```
execv("./a.out", argv);
```

```
if(ret==0) {
```

```
    k=execl("/bin/date", "date", NULL);
```

```
}
```

```
execlp("date", "date", "+%D", NULL);
```

exit is a system call in kernel space

_exit is wrapper for exit system call

exit is a library call in userspace

TODO:- exit vs _exit (exit - lib call, _exit - sys call wrapper)

```
for(i=1;i<=10;i++)  
    printf("hello:%d\t", i);           //no \n, skipping flush  
exit(0);    //_exit(0);                //??
```

Zombie Process / Zombie State.

exit/_exit:- release all resources except PCB and process table entry
zombie state

parent waitpid:- clean up PCB and process table entry
terminated state

process table is of finite length

ulimit -a
ulimit -u

ps
ps aux
ps -el
ps -e -o pid, ppid, stat, cmd

pstree
pstree -np

top # q - quit

Threads:-

Multitasking

Multithreading

- * Concurrent execution
- * Resource Sharing (same address space)
- * increased utilization of CPU

Threads will share all resources, except stack
Every thread maintains private stack

fork vs thread create

==> independent vs shared resources

==> thread creation is faster than fork (no duplication of resources)

Threads are known as Light Weight Process (LWP)

Concurrency Examples (Real Applications)

- * Office Suite
- * browser
- * streaming media player
- * concurrent server
- * parallel sum of large array

Threads -- independent scheduling attributes & context
-- common resource attributes (memory, fs, i/o)

Thread Models:-

- * User level threads, managed by userspace library (thread manager)
kernel is not aware of threads
many to one mapping
e.g. Traditional UNIX
- * Kernel level threads, directly managed by Kernel
one to one mapping
e.g. Linux, Modern OS

ps -eL -o pid,ppid,lwp,nlwp,stat,cmd

Every process runs as single thread initially (LWP)

POSIX Thread Library

[not system calls]

pthread_create

pthread_exit

pthread_join

pthread_equals

pthread_self

prototyped in pthread.h

defined in pthread libs
(libpthread.a or libpthread.so)

/lib, /usr/lib etc

```
void* do_work1(void* pv) {  
    printf("Thread A -- Welcome\n");  
    for(i=1;i<=max;i++)  
        printf("A--%d\n",i);  
}  
//write do_work2 similar to do_work1
```

```
pthread_t pt1;  
pthread_t pt2;  
pthread_create(&pt1, NULL, do_work1, NULL);  
pthread_create(&pt2, NULL, do_work2, NULL);  
----
```

gcc psimple.c -lpthread

exit(0) ==> release all resources
 ==> all threads will terminate

pthread_exit ==> only current thread will terminate
 ==> no release of resources

pthread_join:-

- * block till the completion of specified thread
- * collect exit status of completed thread (for whom waited)

1st param:- pthread_t variables

2nd param:- to collect exit status

pthread_create:-

1st param:- addr of pthread_t variable

2nd param:- attributes, NULL means default

3rd param:- addr/name of service function (callback)

4th param:- arguments to service functions

```
for(i=1;i<=max;i++)          //max can be 10 or 20
{
    printf("A--%d\n",i);
    sleep(1);                  //usleep
}
```



```
void* task_body(void* pv)
```

```
{
```

```
    char* ps=pv;
```

```
    -----
```

```
}
```

```
pthread_create(&pt1,NULL,task_body,"A1");
```

```
pthread_create(&pt2,NULL,task_body,"B2");
```

```
pthread_create(&pt3,NULL,task_body,"C3");
```

```
-----
```

```
int k;
```

```
pthread_t ptarr[n];
```

```
for(i=0;i<n;i++)
```

```
{
```

```
    k=100+i;
```

```
    pthread_create(&ptarr[i],NULL,task_body,(void*)k);
```

```
}
```

```
void* do_service(void* pv) {
```

```
    int k = (int)pv;
```

```
}
```


Task Driven Concurrency
Data Driven Concurrency

Signal Handling:-

Signals -

Asynchronous event .. from one process to other
(software level, kernel level)

one process will generate/trigger signal (sender)
signal will reach other process (target)



Most default handlers cause abnormal termination of a process

familair signal generations.

ctrl + C ==> interrupting process (abnormal term) ==> SIGINT

ctrl + \ ==> SIGQUIT

ctrl + z ==> SIGTSTP

kill <pid> ==> SIGTERM, soft kill

kill -9 <pid> ==> SIGKILL, sure kill

child exit ==> SIGCHLD (to parent)

some other :- SIGSTOP, SIGCONT, SIGFPE, SIGSEGV, SIGALRM

kill -l # list of signals

kill -INT <pid>

kill -2 <pid>

pkill -INT <pname>

pkill -2 <pname>

killall

pkill(pid, signo); # system call

pkill(0, signo); # sending signal to same process

foreground process ==> stdin, stdout focus to user
background process ==> detached from stdin
==> instead of stdout, logging techniques

./a.out # f/g
./a.out& # b/g

command& # run in b/g, no stdin, logging instead of stdout/stderr

jobs ==> list of background and suspended processes

fg ==> resume suspended or b/g process in f/g
bg ==> resume suspended process in b/g

fg <job-id>
bg <job-id>
if no job-id, recent job will be chosen

echo \$? # exit status of recent command

signal API (system call) -- register custom handler
pause -- block until any signal arrives
raise -- send to signal to self
alarm -- sending signal after specified no.of seconds (SIGALRM)

| | | | |
|--------|-----|----------------------|----------------------|
| In PCB | ==> | signal pending field | (collection of bits) |
| | ==> | signal mask field | (collection of bits) |

| | | |
|---------------|-----|---|
| mask bit 0 | ==> | signal will be delivered and handled |
| mask bit 1 | ==> | signal can't be delivered (ignored/masked) |
| pending bit 0 | ==> | such signal not arrived |
| pending bit 1 | ==> | signal will be handled (default/custom handler) |

| | |
|--------------------------|--|
| signal(SIGxxx, handler); | //replace default handler with //custom handler |
|--------------------------|--|

| | |
|------------------|--|
| SIGKILL, SIGSTOP | ==> non maskable signals |
| | ==> custom handler not applicable |
| | ==> only default handler (sure kill, sure suspend) |

Modern signal APIs:-

- * sigaction
- * sigsuspend
- * sigprocmask
- * sigwait
- * kill

```
void handler_for_alarm(int signo) {  
    //print current time  
    //Hints:- time system call, ctime lib API  
    alarm(1);  
}
```

```
signal(SIGALRM, handler_for_alarm);  
alarm(5);  
while(1)  
    pause();
```

Scheduling

TODO:-

- * post read of Threads, Signals + hands-on, examples
- * pre-read on scheduling
- * pre-read on IPC, Semaphores
- * Start Analyzing Assignment-1