

Examples:-

- * race conds , val++, val--
- * prevent race conds using mutex
- * prevent concurrent execution of for loop, using mutex
- * using unnamed semaphore for race conds, concurrent examples

- * named semaphore

producer - consumer scenario

- * a process/thread will add data -- producer
 - * a process/thread will remove data -- consumer
 - * common buffer / data source
-
- * either producer or consumer, only can access common data at a time (shared resource) -- mutual exclusion - R1
 - * consumer should block if buffer empty - R2
 - * producer should block if buffer full - R3

sm	-- semaphore, init to 1 / mutex	-- mutual exclusion
s1	-- semaphore, init to 0	- R2 (cons block if buffer is emp)
s2	-- semaphore, ival? size of buffer	- R3 (prod block if buffer is full)

Producer

Consumer

```
while(1) {
```

```
  lock(s2)
```

```
  lock(sm)
```

```
  //add data,e.g. push
```

```
  unlock(sm)
```

```
  unlock(s1)
```

```
}
```

```
while(1) {
```

```
  lock(s1)
```

```
  lock(sm)
```

```
  //remove data, e.g. pop
```

```
  unlock(sm)
```

```
  unlock(s2)
```

```
}
```

Bounded Buffer Problem

Unbounded Buffer Problem -- producer never blocks

s2 val -- free slots in buffer

s1 val -- filled elements in buffer

Producer

```
while(1) {  
    lock(s2)  
    lock(sm)  
    //.....  
    unlock(sm)  
    unlock(s1)  
}
```

Consumer

```
while(1) {  
    lock(sm)  
    lock(s1)    //?? : dead lock  
    //.....  
}
```

Avoid - Apply mutual exclusion, after resolving dependencies

Spinlock

Condition variables

Spinlock vs Semaphore

==> if cond is not met, Semaphore will block the process

==> if cond is not met, Spinlock will hold proces in a busy loop (spinning)

Spinlock:- flag=0;

P1

P2

lock:-

lock(s1)

lock(s1)

while(flag);
flag=1;

//critical

//critical

unlock:-

flag=0;

unlock(s1)

unlock(s1)

Atomic instructions:-
XCHG or SWP

If critical section is small + SMP
==> spinlocks
if critical section is big (or) no SMP
==> semaphores

flag=0
r1 <--1
while(XCHG(r1,flag)); //perfect soln

H/w level atomic execution
* disable switching
* data bus locking to avoid
variable access by other CPUs

File System Management:-

File System? Organization of files & dirs, content in files

create, modify, delete, rename, copy, move files

File Systems Impl -->

Real/Physical Systems:-

FAT32, NTFS, ext2, ext3, ext4, XFS, ZFS, HFS, reiserfs, cramfs

Pseudo Filesystems:-

procfs, sysfs, devfs, debugfs, tmpfs

{ /proc /sys /dev /sys/kernel/debug, /var/tmp }

File System -- Logical

Disk/Storage -- Physical

(SATA, IDE/PATA, SSD, USB Storage)

FS : Blocks, Disk: Sectors

Virtual File System(VFS) -- gateway for all other file systems(common bridge)

uspsace read --> vfs read --> fat32 read --> sata read
--> xfs read --> ssd read

format ?? making new file system
mounting

5 parts on disk:- /dev/sda1, /dev/sda5, /dev/sda6, /dev/sda7, /dev/sda8
 C: D: E: F: G:

Pen drive with no part:- /dev/sdb

(or) pen drive with two parts:- /dev/sdb1, /dev/sdb2

/ -- root partition (sda7)

/dev/sda1 -- /mnt/c mount /dev/sda1 /mnt/c

/dev/sda5 -- /mnt/d

/dev/sda6 -- /mnt/e

/dev/sda8 -- /mnt/g

/dev/sdb -- /mnt/usb mount /dev/sdb /mnt/usb


umount /mnt/c

umount /mnt/c

Automount ==> /media/xxx


FS Layout:-

R1 -- boot sector/block
R2 -- super block
R3 -- inode block
R4 -- data blocks



Block size, e.g. 4 KB
File of 40 KB

i-node:-
data structure/descriptor
holding file attributes
file descriptor



ls -i ==> inode numbers

ls -li ==> inode numbers + file details

File Attributes Stored in i-node:-

- * block info (of file)
- * permissions
- * type of file (regular, dir, symlink, pipe/fifo, char spl, block spl)
- * user, group
- * file size
- * timestamps - mtime, atime, ctime

ls -l command (or) lstat example

file name vs inode number mapping ==> directory

/home/rajesh/os/test/sample.txt

inode no. of / ==> 2
ls -lid /

Self Study:-

hard & soft links, ln command

lstat system call

chmod, chown, umask -- permissions & ownership

file handling -- open, read, write, close

write:-

```
fd=open("sample.txt",O_WRONLY);           //O_WRONLY|O_CREAT, 0666
char msg[]="Hello Linux";
nbytes=write(fd, msg, len);
close(fd);
```

//alpha.txt:- ABCD..XYZ0123..9 //abcd...xyz

```
fd=open("sample.txt",O_RDONLY);
char buf[128];
int maxlen=128;
nbytes=read(fd,buf,maxlen);      //nbytes:36
//print buf
close(fd);
```

```
int buflen=10;
read(fd, buf, buflen);      //AB..IJ,    offset:10, ret:10
read(fd, buf, buflen);      //KL..ST,    offset:20      ret:10
read(fd, buf, buflen);      //UV...0123, offset:30      ret:10
read(fd, buf, buflen);      //456789, offset:36 ret:6
```


alpha.txt ==> 36 bytes

```
fd = open("alpha.txt", O_RDONLY);
```

nbytes=read(fd,buf, 10);	//off:10, AB...IJ	SEEK_SET
lseek(fd, 15, SEEK_SET);	//off:15	SEEK_CUR
nbytes=read(fd,buf,6);	//buf:PQRSTU, off:21	SEEK_END
lseek(fd, -8, SEEK_CUR);	//off:13	
nbytes=read(fd,buf,5);	//buf:NOPQR, off:18	
lseek(fd,-12, SEEK_END);	//off:24 (36-12)	

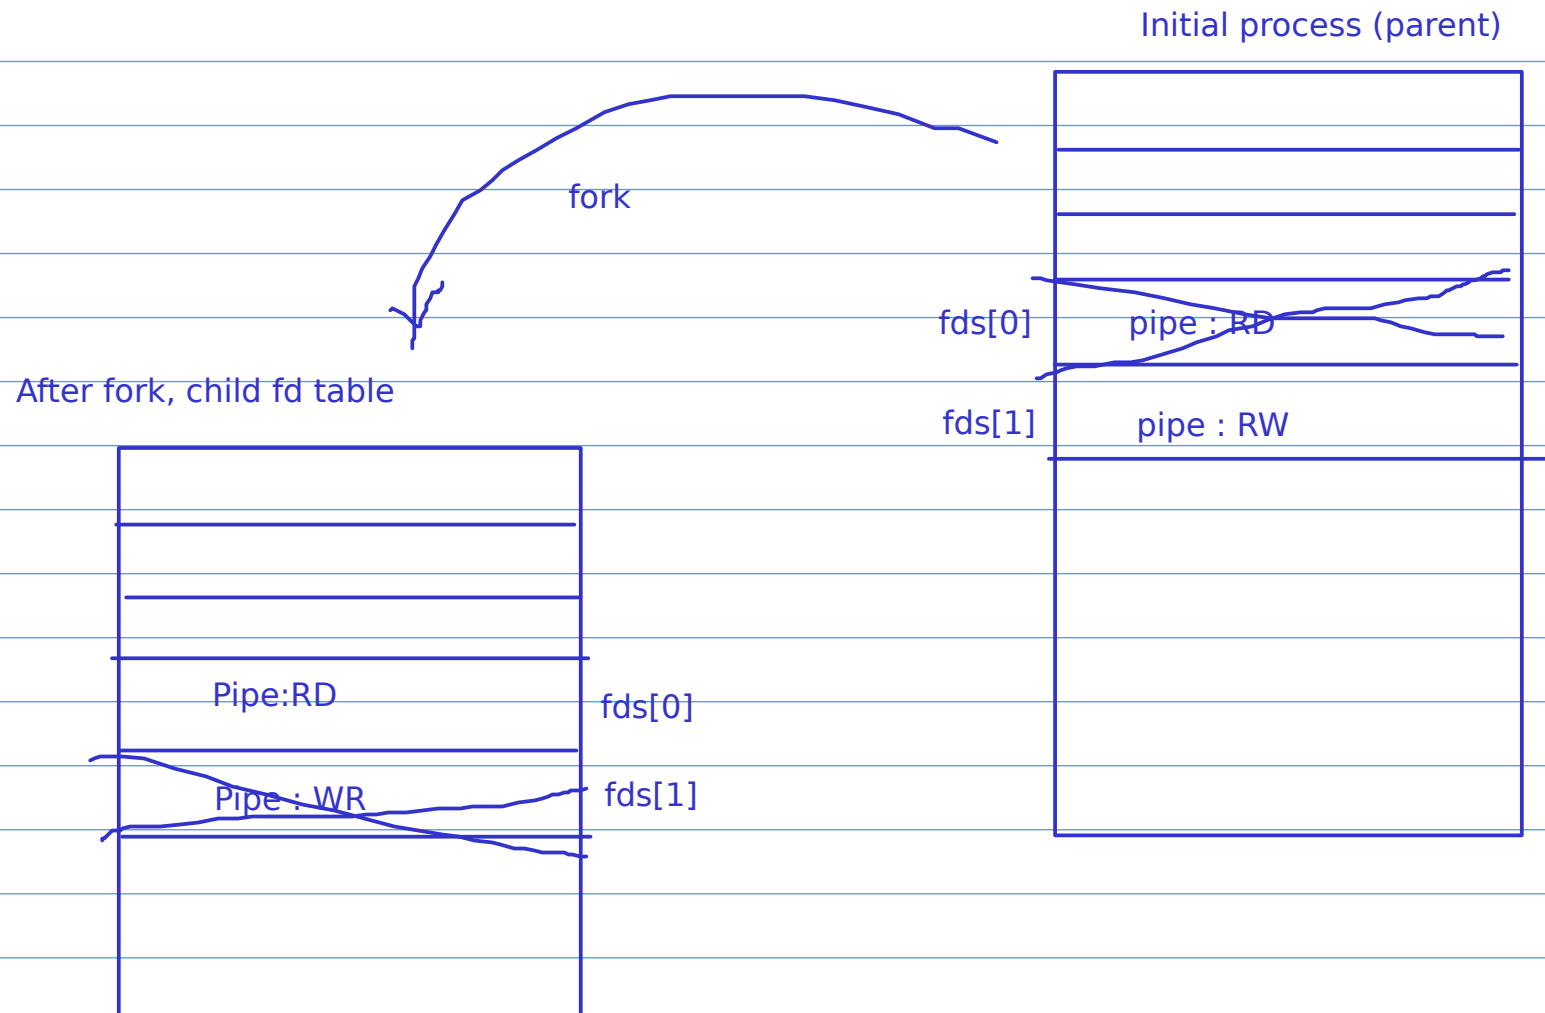
pipe / fifo:-

- * used for IPC (data exchange)
- * pseudo file
- * one process can write (one end), othe process can read(other end)
- * stream of data, FIFO order
- * reading an empty pipe ==> block process

unnamed pipe (pipe)	-- no file name, memory inode exists
	-- applicable for related processes only (parent-child)

named pipe (FIFOs)	-- file name applicable, disk inode exists
	-- but data blocks in memory
	-- applicable for any two processes

```
int fds[2];  
ret=pipe(fds); //fd[0] -- read, fds[1] -- write
```



TODO:-
Named Pipes (FIFOs)

Further Topics:-

Message Queues
Virtual Memory
Shared Memory
Scheduling