

Further Topics:-

Hint:- futex

Named Pipes (FIFOs)
Message Queues
Scheduling

Virtual Memory
Shared Memory

Named Pipes/FIFOs

mkfifo command
mkfifo lib API
mknod system call

mkfifo /tmp/p1
ls -l /tmp/p1

#In terminal-1
cat /tmp/p1

#In terminal-2
echo "abcdxyz" > /tmp/p1

mkfifo /tmp/p2

writer.c:-
fd=open("/tmp/p2", O_WRONLY);
//err handling
write(fd,msg,len);
//err handling
close(fd);

reader.c:-
fd=open("/tmp/p2", O_RDONLY);
//err handling
read(fd,buf,maxlen);
//print buf
//err handling
close(fd);

file <xxxxx> # based on header

IPC:- Synchronization, Data Exchange

Sys V Semaphores:-

semget, semop, semctl

POSIX Semaphores:-

sem_init, sem_open, sem_destroy

sem_wait, sem_post

Sys V Message Queues:-

msgget, msgsnd, msgrcv, msgctl

POSIX Message Queues:-

mq_open, mq_send, mq_receive, mq_unlink

Message Queues:-

Purpose:- Data Exchange

Semaphores

Message Queues

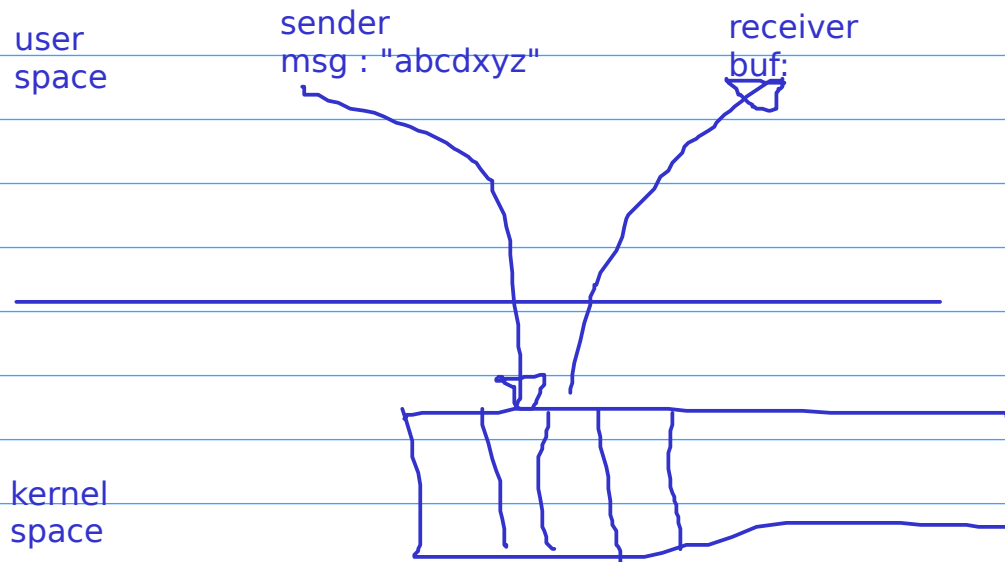
Shared Memory

|

Sys V Variant

POSIX Variant

POSIX



Internal sync support
Typically FIFO
Discrete messages:-

send m1(20)
send m2(10)

recv with 40 bytes
--> m1 : 20 bytes
recv with 15 bytes
--> m2 : 10 bytes
recv
--> block

Steps:-

```
int mqfd;  
mqfd = mq_open("/mqsimple", O_RDWR);  
mqfd = mq_open("/mqsimple", O_RDWR|O_CREAT,0666,NULL);
```

```
sender:-  
char msg[] = "Hello POSIX";  
int len=strlen(msg);  
mq_send(mqfd, msg, len, 5);
```

receiver:-

```
char buf[MAXLEN];  
int maxlen=8192;  
int prio=5;  
mq_receive(mqfd, buf, maxlen, &prio);
```

close:-

```
mq_close(mqfd);
```

Default msg len : 8192

During receive - maxlen \geq msg len
//through attr, can override default msg len

```
mqd_t mqid;
```

```
struct mq_attr attr;
```

```
attr.mq_msgsize=512;
```

```
attr.mq_maxmsg=10;
```

```
mqid=mq_open("/mque2",O_WRONLY|O_CREAT,0666,&attr);
```

linker flag:-

-lrt

```
gcc sender.c -lrt -o sender
```

```
ls /dev/mqueue
```

```
cat /dev/mqueue/mqsimple
```

```
man mq_overview
```

```
mq_setattr
```

```
mq_getattr
```

```
mq_unlink
```

Scheduling:-

When do CPU becomes free? When switching required?

- * task completion A
- * blocking call e.g sleep, read, mq_receive, sem_wait B
- * timeout C
- * interrupt arrives, end of ISR D

A & B only -- non preemptive

C & D also -- preemptive

Note:- process never sleep/block in userspace
blocking always occur in kernel space (system call)

Classical Policies:-

FCFS/FIFO

Priority Based

Round Robin (Time slice)

Short Job First

O(1) complexity expected -- idle scheduling

Q1 -- telephonic	: Priority (real time)	FCFS
Q2 -- user interactive	: RR (fairness)	RR
Q3 -- background/maintenance	: FCFS (simplicity)	PRI + FCFS
	: RR with larger quantum	PRI + RR

processes may move across queues

- explicit change in priority/process
- aging
- downgrade (lot of CPU usage)

Priority policies :- different queues for each priority (goal $O(1)$ complexity)

Scheduling in Linux:-

* Real time policies : 100 priority levels : 0 - 99 (0-min, 99-max)

SCHED_FIFO : Priority + Fifo

SCHED_RR : Priority + RR

* Non real time policies

SCHED_OTHER : Time sharing algo

: Modified Round Robin

ps -e -o pid,ppid,stat,policy,pri,cmd

sudo chrt -f 5 ./a.out

sudo chrt -r 15 ./a.out

True/Complete Fairness:-

- * Every process should be scheduled atleast once, before any other process goes for second turn

Limitation of Round Robin:-

Time slice - 2 units

	Time	CPU	Queue
p1 - 0, 4 units	0	p1	empty
p2 - 1, 6 units	1	p1	p2, p3
p3 - 1	1.9	p1	p2,p3,p4
p4 - 1.9	2	p2	p3,p4,p1
p5 - 3	3	p2	p3,p4,p1,p5

Time sharing:-

- * Two queues -- active, expired
- * Preempted process (timeout) will go back to expired
- * New processes are added to active queue
- * Scheduler always pickup from Active queue
- * When active Q is empty, swapping occurs b'n active & expired
- * Priority is applicable , 40 levels
- * Pair of active & expires for each prio level

* Dynamic time slice -- proportional to priority

Nice vals -- controls priority in TS Algo (inverse relation)

Range : -20 to +19
default nice value : 0

+6, less prior than 0
-9, more prior than 0

-20 : Highest, 0 : default, +19: lowest

ni	:	0	5	-6	-20	+19	(-20 to +19 scale)
----	---	---	---	----	-----	-----	--------------------

ps o/p	:	19	14	25	39	0	(0 to 39 scale, 0-low, 39-high)
--------	---	----	----	----	----	---	---------------------------------

nice -n 5 ./a.out

nice -n -6 ./a.out

#renice to change ni value of existing

lowering nice vals, -ve ni values are not allowed for normal user
(only for admin/root user)

APIs:- sched_setscheduler,sched_setparam

SMT or Hyperthreading:-

Switching b'n two threads at h/w level

Core i3

Core i5

Core i7

nproc

cat /proc/cpuinfo

phy

threads

2

2

2

4 *

4

4

4

8 *

AMP : asymmetric, co-processor

e.g. GPU, ARM + DSP

Scheduling in SMP

* SMP : multiple CPU cores, identical

SMT

Symmetric Multi Threading

* Common ready Q or private ready Qs ?

e.g. Hyperthreading

* Modern kernels (e.g. Linux kernel > v2.6)

==> private ready Q

* CPU Affinity (No Migration), reason - private cache entries

Migration -- discard CPU cache, rebuild on other CPU

0x1 -- CPU0

0x2 -- CPU1

Load Balancing Techniques, e.g. Pull Migration, Push Migration

0x4 -- CPU2

0x8 -- CPU3

taskset 0x1 ./receiver

taskset -p 0x4 <pid>

0x6 -- 1 or 2

pthread_setaffinity_np, pthread_attr_setaffinity_np

Memory Management / Virtual Memory:-

Myths / Facts:-

* 1 GB RAM + 2 GB swap area	# swap area -- from disk
virtual memory size? 3 GB	

1 GB	==>	2^{30}
4 GB	==>	2^{32}

Address < 1GB ==> RAM
Address > 1GB, but < 3 GB ==> Not Applicable

Whichever address, CPU will access RAM only (not swap area for active load/store)

Processes (pages) in swap area are not under active execution

Swap area -- backing store

4 GB RAM + 0 swap area ==> 4 GB of Virtual Mem

Virtual Memory Management:-

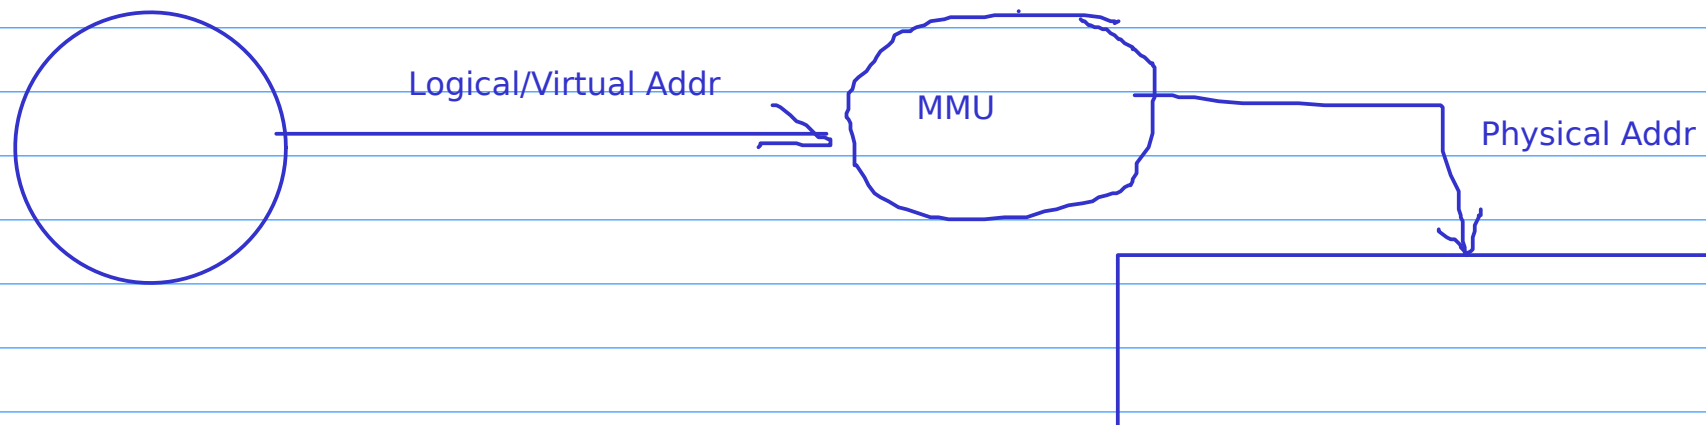
- * Allocate/Deallocate memory for processes (stack, heap, code, idata, udata)

challenges:-

- * How to handle scalable memory
- * holes in memory - how to manage with non contiguous memory
- * Expand physical memory (+swap area)
- * Protection b'n processes (Addresss space)

Memory Mapping:-

compiler & linker generates logical addresses, which are issued by CPU



Pages, Paging:-

- * Physical memory is divided into equal size partitions -- Frames
- * Logical memory is also divided into same size parts -- Pages
- * Page/Frame size is fixed by h/w design

* Pages will be contiguous, but not frames (not necessary)

* Mapping of page no.s and frame no.s --> Page Table

* Page Table -- collection of Page Table Entries (PTEs)

In x86
page size
4 KB

0-4095 - p0 - 15

4096-8191 - p1 - 25

8192-12287 - p2 - 20

12287-16383 - p3 - 12

Logical addr: 4200 ==> p1, offset: 104

Logical addr: 4500 ==> p1, offset: 404

Logical addr: 8400 ==> p2, offset: 208